

Diogo Ribeiro

*ESMAD - Escola Superior de Média Arte e Design  
Lead Data Scientist, Mysense.ai*

✉ dfr@esmad.ipp.pt   >ID 0009-0001-2022-7072

November 16, 2025

# Outline

- 1 Introduction
- 2 Neural Network Basics
- 3 Backpropagation and Training
- 4 Regularization
- 5 Convolutional Neural Networks (CNN)
- 6 Recurrent Neural Networks (RNN)
- 7 Transformer Architecture
- 8 Practical Training Tips
- 9 Conclusion



# What is Deep Learning?

## Deep Learning

A class of machine learning algorithms that use **multiple layers** of non-linear transformations to learn hierarchical representations from data.

### Key characteristics:

- **Depth:** Many layers (hence "deep")
- **End-to-end learning:** Learn features automatically
- **Hierarchical representations:** Low-level → High-level
- **Scalability:** Performance improves with more data

### Why now?

- Big Data: Massive datasets available
- Compute: GPUs, TPUs enable fast training
- Algorithms: Better architectures and training techniques

# The Biological Inspiration

## Artificial neurons inspired by biological neurons:

### Biological Neuron:

- Dendrites: Receive signals
- Cell body: Processes signals
- Axon: Sends output
- Synapses: Connect neurons

### Artificial Neuron:

- Inputs:  $x_1, \dots, x_n$
- Weights:  $w_1, \dots, w_n$
- Aggregation:  $\sum w_i x_i + b$
- Activation:  $f(\cdot)$

### Important

Modern deep learning has moved beyond simple biological analogy!

- Backpropagation doesn't occur in brain
- Architectures are engineering solutions
- Biological plausibility not the goal



# The Perceptron

## Perceptron (Rosenblatt, 1958)

$$y = f \left( \sum_{i=1}^n w_i x_i + b \right) = f(\mathbf{w}^T \mathbf{x} + b) \quad (1)$$

where  $f$  is a step function:  $f(z) = \begin{cases} 1 & z \geq 0 \\ 0 & z < 0 \end{cases}$

### Learning rule:

$$w_i \leftarrow w_i + \eta(y_{\text{true}} - y_{\text{pred}})x_i \quad (2)$$

where  $\eta$  is the learning rate.

### Limitations (Minsky & Papert, 1969):

- Can only learn linearly separable functions

# Multi-Layer Perceptron (MLP)

**Solution:** Add hidden layers!

## Feed-Forward Neural Network

**Layer  $l$ :**

$$\mathbf{z}^{(l)} = \mathbf{W}^{(l)} \mathbf{a}^{(l-1)} + \mathbf{b}^{(l)} \quad (3)$$

$$\mathbf{a}^{(l)} = f(\mathbf{z}^{(l)}) \quad (4)$$

where:

- $\mathbf{a}^{(l)}$ : Activations (outputs) of layer  $l$
- $\mathbf{W}^{(l)}$ : Weight matrix
- $\mathbf{b}^{(l)}$ : Bias vector
- $f$ : Activation function

# Activation Functions

**Non-linearity is crucial! Without it, network is just linear regression.**

## 1. Sigmoid:

$$\sigma(z) = \frac{1}{1 + e^{-z}} \quad (5)$$

- Output:  $(0, 1)$
- Smooth gradient
- Problem: Vanishing gradients

## 2. Tanh:

$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}} \quad (6)$$

- Output:  $(-1, 1)$

## 3. ReLU:

$$\text{ReLU}(z) = \max(0, z) \quad (7)$$

- Most popular!
- No vanishing gradients
- Fast computation
- Problem: "Dead neurons"

## 4. Variants:

- **Leaky ReLU:**  $\max(0.01z, z)$
- **ELU:** Smooth for  $z < 0$



# The Learning Problem

**Goal:** Minimize loss function  $\mathcal{L}(\theta)$  over parameters  $\theta$ .

**Common loss functions:**

- **Mean Squared Error (Regression):**

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (8)$$

- **Cross-Entropy (Classification):**

$$\mathcal{L} = -\frac{1}{n} \sum_{i=1}^n \sum_{c=1}^C y_{ic} \log(\hat{y}_{ic}) \quad (9)$$

- **Binary Cross-Entropy:**

$$\mathcal{L} = -\frac{1}{n} \sum_{i=1}^n [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)] \quad (10)$$

# Backpropagation

**Efficiently compute gradients using chain rule!**

## Chain Rule

For composite function  $f(g(x))$ :

$$\frac{df}{dx} = \frac{df}{dg} \cdot \frac{dg}{dx} \quad (11)$$

**For neural networks:**

**Forward pass:** Compute activations layer by layer

$$\mathbf{a}^{(l)} = f(\mathbf{W}^{(l)} \mathbf{a}^{(l-1)} + \mathbf{b}^{(l)}) \quad (12)$$

**Backward pass:** Compute gradients layer by layer

$$\delta^{(l)} = (\mathbf{W}^{(l+1)})^T \delta^{(l+1)} \odot f'(\mathbf{z}^{(l)}) \quad (13)$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{a}^{(l)}} = \delta^{(l)} \mathbf{a}^{(l-1) \top T} \quad (14)$$

## Update rule:

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} \mathcal{L}(\theta_t) \quad (15)$$

## Variants:

- **Batch Gradient Descent:** Use full dataset
  - Accurate gradients
  - Slow for large datasets
- **Stochastic Gradient Descent (SGD):** One sample at a time
  - Fast updates
  - Noisy gradients
- **Mini-Batch SGD:** Batches of  $B$  samples
  - Balance between accuracy and speed
  - Most commonly used
  - Typical batch sizes: 32, 64, 128, 256

# Advanced Optimizers

**Momentum (Polyak, 1964):**

$$\mathbf{v}_t = \beta \mathbf{v}_{t-1} + \nabla_{\theta} \mathcal{L}(\theta_t) \quad (16)$$

$$\theta_{t+1} = \theta_t - \eta \mathbf{v}_t \quad (17)$$

Accumulates velocity, smooths updates.

**Adam (Kingma & Ba, 2015):** Most popular!

$$\mathbf{m}_t = \beta_1 \mathbf{m}_{t-1} + (1 - \beta_1) \nabla_{\theta} \mathcal{L} \quad (18)$$

$$\mathbf{v}_t = \beta_2 \mathbf{v}_{t-1} + (1 - \beta_2) (\nabla_{\theta} \mathcal{L})^2 \quad (19)$$

$$\hat{\mathbf{m}}_t = \mathbf{m}_t / (1 - \beta_1^t), \quad \hat{\mathbf{v}}_t = \mathbf{v}_t / (1 - \beta_2^t) \quad (20)$$

$$\theta_{t+1} = \theta_t - \eta \frac{\hat{\mathbf{m}}_t}{\sqrt{\hat{\mathbf{v}}_t} + \epsilon} \quad (21)$$

- Adaptive learning rates per parameter
- Default:  $\beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}$



# Why Regularization?

## The Overfitting Problem

Deep networks have millions of parameters → Can memorize training data!

### Consequences:

- Perfect training accuracy
- Poor generalization to test data
- Unreliable predictions

### Regularization techniques:

1. L1/L2 weight penalties
2. Dropout
3. Batch normalization
4. Data augmentation
5. Early stopping

# L1 and L2 Regularization

Add penalty term to loss function:

L2 (Ridge):

$$\mathcal{L}_{\text{reg}} = \mathcal{L} + \lambda \sum_i w_i^2 \quad (22)$$

- Shrinks weights towards zero
- Smooth, differentiable
- Most common in deep learning

L1 (Lasso):

$$\mathcal{L}_{\text{reg}} = \mathcal{L} + \lambda \sum_i |w_i| \quad (23)$$

- Encourages sparsity
- Some weights  $\rightarrow$  exactly zero
- Acts as feature selection

## Dropout (Srivastava et al., 2014)

During training, randomly set a fraction  $p$  of neurons to zero.

**Training:**

$$\mathbf{a}^{(l)} = f(\mathbf{z}^{(l)}) \odot \mathbf{m}, \quad \mathbf{m} \sim \text{Bernoulli}(1 - p) \quad (24)$$

**Testing:** Use all neurons, scale by  $(1 - p)$

## Why does it work?

- Ensemble effect: Training many "thinned" networks
- Prevents co-adaptation of neurons
- Forces redundant representations

**Typical values:**  $p = 0.5$  for hidden layers,  $p = 0.1$  for input

# Batch Normalization

## Batch Normalization (Ioffe & Szegedy, 2015)

Normalize activations across mini-batch:

$$\mu_B = \frac{1}{B} \sum_{i=1}^B \mathbf{z}_i \quad (25)$$

$$\sigma_B^2 = \frac{1}{B} \sum_{i=1}^B (\mathbf{z}_i - \mu_B)^2 \quad (26)$$

$$\hat{\mathbf{z}}_i = \frac{\mathbf{z}_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \quad (27)$$

$$\mathbf{y}_i = \gamma \hat{\mathbf{z}}_i + \beta \quad (28)$$

where  $\gamma, \beta$  are learned parameters.



# Why CNNs for Images?

## Fully-connected networks don't scale:

- $224 \times 224$  RGB image: 150,528 inputs
- 1000 hidden units: 150M parameters (just first layer!)
- Ignores spatial structure

## Key principles of CNNs:

1. **Local connectivity:** Neurons connect to small regions
2. **Parameter sharing:** Same weights applied everywhere
3. **Translation equivariance:** Shift input  $\rightarrow$  Shift output

## Biological Inspiration

Based on Hubel & Wiesel's (1962) work on cat visual cortex:

- Simple cells: Detect edges
- Complex cells: Invariant to position

# Convolution Operation

## 2D Convolution

$$(I * K)[i, j] = \sum_m \sum_n I[i + m, j + n] \cdot K[m, n] \quad (29)$$

where  $I$  is input image,  $K$  is kernel (filter).

### Hyperparameters:

- **Kernel size:** Usually  $3 \times 3$  or  $5 \times 5$
- **Stride:** Step size (usually 1 or 2)
- **Padding:** "same" or "valid"
- **Number of filters:** Determines output depth

### Output size:

$$O = \left\lfloor \frac{I + 2P - K}{S} \right\rfloor + 1 \quad (30)$$

# CNN Architecture Components

## Typical CNN structure:

### 1. Convolutional Layers:

- Apply learned filters
- Extract local features
- Multiple filters → Multiple feature maps

### 2. Pooling Layers:

- Reduce spatial dimensions
- **Max pooling:** Take maximum in each region
- **Average pooling:** Take average
- Provides translation invariance

### 3. Fully-Connected Layers:

- At the end of network
- Combine features for final prediction

**Example:** Conv → ReLU → Pool → Conv → ReLU → Pool → FC → Softmax

# Classic CNN Architectures

## 1. LeNet-5 (LeCun, 1998):

- First successful CNN
- Digit recognition (MNIST)
- 7 layers, 60K parameters

## 2. AlexNet (Krizhevsky et al., 2012):

- ImageNet breakthrough (top-5 error: 15.3%)
- 8 layers, 60M parameters
- ReLU, dropout, data augmentation
- Trained on GPUs

## 3. VGGNet (Simonyan & Zisserman, 2014):

- Very deep (16-19 layers)
- Small  $3 \times 3$  filters throughout
- 138M parameters

## 4. ResNet (He et al., 2015):

- Residual connections enable very deep networks (50-152 layers)
- Skip connections (Residual blocks)



# Sequential Data

**Many problems involve sequences:**

- Text: Words in sentences
- Speech: Audio over time
- Time series: Stock prices, weather
- Video: Frames over time

**Challenge:** Variable-length inputs and outputs

**RNN solution:** Maintain hidden state that captures history

$$\mathbf{h}_t = f(\mathbf{W}_{hh}\mathbf{h}_{t-1} + \mathbf{W}_{xh}\mathbf{x}_t + \mathbf{b}_h) \quad (31)$$

**Key idea:** Share parameters across time steps!

## Different input/output configurations:

1. **One-to-One:** Standard neural network
2. **One-to-Many:** Image captioning (image → sequence)
3. **Many-to-One:** Sentiment analysis (sequence → label)
4. **Many-to-Many (same length):** Video classification
5. **Many-to-Many (different length):** Machine translation

## Challenges:

- **Vanishing gradients:** Hard to learn long-term dependencies
- **Exploding gradients:** Unstable training
- **Solution:** LSTM and GRU cells

# LSTM (Long Short-Term Memory)

Hochreiter & Schmidhuber (1997) - Addresses vanishing gradients

Key components:

- **Cell state  $c_t$ :** Long-term memory highway
- **Gates:** Control information flow

LSTM equations:

$$f_t = \sigma(\mathbf{W}_f[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_f) \quad (\text{forget gate}) \quad (32)$$

$$i_t = \sigma(\mathbf{W}_i[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_i) \quad (\text{input gate}) \quad (33)$$

$$\tilde{c}_t = \tanh(\mathbf{W}_c[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_c) \quad (\text{candidate}) \quad (34)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t \quad (\text{cell state}) \quad (35)$$

$$o_t = \sigma(\mathbf{W}_o[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_o) \quad (\text{output gate}) \quad (36)$$

$$h_t = o_t \odot \tanh(c_t) \quad (\text{hidden state}) \quad (37)$$



# Attention is All You Need

Vaswani et al. (2017) - Revolutionary architecture

Problems with RNNs:

- Sequential computation (can't parallelize)
- Long-range dependencies still challenging
- Slow training

Transformer solution:

- **Self-attention:** Every position attends to all positions
- **Parallelizable:** No sequential dependency
- **Positional encoding:** Add position information

## Impact

Transformers now dominate NLP (BERT, GPT) and expanding to vision (ViT), audio, and more!

## Scaled Dot-Product Attention

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}}\right) \mathbf{V} \quad (38)$$

where:

- $\mathbf{Q}$ : Query matrix (what we're looking for)
- $\mathbf{K}$ : Key matrix (what's available)
- $\mathbf{V}$ : Value matrix (actual content)
- $d_k$ : Dimension of keys (for scaling)

## Multi-Head Attention:

- Run attention mechanism multiple times in parallel
- Each "head" learns different patterns
- Concatenate and sum the outputs

# Transformer Architecture

## Encoder-Decoder structure:

### Encoder (left side):

1. Multi-head self-attention
2. Add & Normalize
3. Feed-forward network
4. Add & Normalize
5. (Repeat  $N$  times)

### Decoder (right side):

1. Masked multi-head self-attention
2. Add & Normalize
3. Cross-attention to encoder
4. Add & Normalize
5. Feed-forward network
6. Add & Normalize
7. (Repeat  $N$  times)

# Transformer Variants

**Major models based on Transformers:**

**BERT (Devlin et al., 2018):**

- Encoder-only
- Bidirectional context
- Pre-training: Masked language modeling
- Fine-tuning for downstream tasks

**GPT (Radford et al., 2018-2023):**

- Decoder-only
- Autoregressive generation
- Scaling to 175B+ parameters (GPT-3)
- In-context learning

**Vision Transformer (Dosovitskiy et al., 2020):**

- Apply Transformers to images
- Split image into patches
- Competitive with CNNs on large datasets



# Training Best Practices

## 1. Data Preprocessing:

- Normalize inputs (zero mean, unit variance)
- Data augmentation (images: flip, crop, rotate)
- Handle class imbalance

## 2. Initialization:

- **Xavier/Glorot**: For tanh/sigmoid
- **He initialization**: For ReLU
- Never initialize all weights to zero!

## 3. Learning Rate:

- Start with default values (0.001 for Adam)
- Learning rate schedules (decay, cyclical)
- Warm-up for large batches

## 4. Batch Size:

- Larger batches: Faster, more stable
- Smaller batches: Better generalization
- Trade-off with GPU memory

# Debugging Neural Networks

## Common issues and solutions:

### 1. Loss not decreasing:

- Check learning rate (too high or too low)
- Verify gradient flow (print gradient norms)
- Check data preprocessing

### 2. Loss exploding:

- Reduce learning rate
- Gradient clipping
- Check for bugs in loss computation

### 3. Overfitting:

- Add regularization (dropout, L2)
- More data or data augmentation
- Reduce model capacity

### 4. Underfitting:

- Increase model capacity
- Train longer
- Reduce regularization



# Summary

## Key concepts covered:

- **Foundations:** Neurons, activation functions, backpropagation
- **Optimization:** SGD, Adam, learning rate schedules
- **Regularization:** Dropout, batch normalization, weight decay
- **CNNs:** Convolution, pooling, ResNet
- **RNNs:** LSTM, sequence modeling
- **Transformers:** Attention, BERT, GPT

## The Deep Learning Revolution

- Transforming AI and industry
- Rapid progress in architectures and applications
- Democratization through frameworks (PyTorch, TensorFlow)
- Still many open questions!

# Further Reading

## Books:

- Goodfellow, Bengio, Courville (2016). *Deep Learning*
- Zhang et al. (2023). *Dive into Deep Learning*

## Courses:

- Stanford CS231n (CNNs for Visual Recognition)
- Stanford CS224n (NLP with Deep Learning)
- fast.ai Practical Deep Learning

## Key Papers:

- Krizhevsky et al. (2012). "ImageNet Classification with Deep CNNs" (AlexNet)
- He et al. (2015). "Deep Residual Learning" (ResNet)
- Vaswani et al. (2017). "Attention Is All You Need" (Transformer)

# Acknowledgments

- ESMAD for institutional support
- Mysense.ai for deep learning applications
- Deep learning research community

Generated with **LATEX** Beamer

Theme: ESMAD Professional Academic Style

## Diogo Ribeiro

ESMAD - Escola Superior de Média Arte e Design  
Lead Data Scientist, Mysense.ai

-  dfr@esmad.ipp.pt
-  0009-0001-2022-7072
-  [github.com/diogoribeiro7](https://github.com/diogoribeiro7)
-  [linkedin.com/in/diogoribeiro7](https://linkedin.com/in/diogoribeiro7)

This presentation is part of the academic materials repository.

Available at: <https://github.com/diogoribeiro7/academic-presentations>