

Bayesian Machine Learning

Neural Networks and Gaussian Processes

Diogo Ribeiro
ESMAD – Escola Superior de Média Arte e Design
Lead Data Scientist, Mysense.ai

October 27, 2025

Outline

- 1 Introduction to Bayesian Machine Learning
- 2 Bayesian Neural Networks
- 3 Gaussian Processes
- 4 Advanced Topics and Modern Developments
- 5 Practical Implementation and Software
- 6 Future Directions and Conclusions

Why Bayesian Machine Learning?

Traditional ML Limitations:

- **Point estimates:** Single "best" parameters
- **Overconfidence:** No uncertainty quantification
- **Overfitting:** Limited regularization mechanisms
- **Model selection:** Ad-hoc validation approaches

Bayesian Advantages:

- **Uncertainty quantification:** Principled confidence intervals
- **Automatic regularization:** Prior knowledge integration
- **Model comparison:** Marginal likelihood for selection
- **Sequential learning:** Natural online updates

Bayesian Paradigm

$$p(\boldsymbol{\theta}|\mathcal{D}) = \frac{p(\mathcal{D}|\boldsymbol{\theta})p(\boldsymbol{\theta})}{p(\mathcal{D})} \quad (1)$$

$$\text{Posterior} = \frac{\text{Likelihood} \times \text{Prior}}{\text{Evidence}} \quad (2)$$

Key Insight

Treat parameters as random variables, not fixed unknowns

Bayesian vs Frequentist Perspectives

Aspect	Frequentist	Bayesian
Parameters	Fixed unknown constants	Random variables with distributions
Uncertainty	Confidence intervals (repeated sampling)	Credible intervals (probability statements)
Model Selection	Cross-validation, AIC/BIC	Marginal likelihood, posterior odds
Regularization	L_1/L_2 penalties	Prior distributions
Prediction	Point estimates	Predictive distributions
Computational	Optimization-based	Integration-based (MCMC/VI)

Computational Challenges and Solutions

The Integration Problem:

$$p(\boldsymbol{\theta}|\mathcal{D}) = \frac{p(\mathcal{D}|\boldsymbol{\theta})p(\boldsymbol{\theta})}{\int p(\mathcal{D}|\boldsymbol{\theta})p(\boldsymbol{\theta})d\boldsymbol{\theta}}$$

Challenges:

- High-dimensional integrals
- No closed-form solutions
- Computational complexity
- Scalability to big data

Historical Solutions:

- Conjugate priors
- Laplace approximation

Modern Approaches:

Markov Chain Monte Carlo

- Hamiltonian Monte Carlo
- No-U-Turn Sampler (NUTS)
- Exact sampling (asymptotically)

Variational Inference

- Mean-field approximation
- Normalizing flows
- Automatic differentiation
- Scalable to large datasets

Trade-offs

MCMC • Exact but slow

From Neural Networks to Bayesian Neural Networks

Standard Neural Network:

$$y = f(\mathbf{x}; \mathbf{w}) + \epsilon$$

Learning: Find optimal weights \mathbf{w}^*

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) + \lambda R(\mathbf{w})$$

Prediction: Point estimate

$$p(y^*|\mathbf{x}^*) = \delta(y^* - f(\mathbf{x}^*; \mathbf{w}^*))$$

Issues:

- No uncertainty quantification
- Prone to overconfidence

Bayesian Neural Network:

$$y = f(\mathbf{x}; \mathbf{w}) + \epsilon, \quad \mathbf{w} \sim p(\mathbf{w})$$

Learning: Posterior distribution

$$p(\mathbf{w}|\mathcal{D}) \propto p(\mathcal{D}|\mathbf{w})p(\mathbf{w})$$

Prediction: Averaging over weights

$$p(y^*|\mathbf{x}^*, \mathcal{D}) = \int p(y^*|\mathbf{x}^*, \mathbf{w})p(\mathbf{w}|\mathcal{D})d\mathbf{w}$$

Benefits:

- Uncertainty quantification
- Automatic regularization

BNN Architecture and Prior Specification

Network Architecture:

$$h_1 = \sigma(W_1 \mathbf{x} + b_1) \quad (3)$$

$$h_2 = \sigma(W_2 h_1 + b_2) \quad (4)$$

$$\vdots \quad (5)$$

$$y = W_L h_{L-1} + b_L \quad (6)$$

Prior Distributions:

$$W_{ij}^{(l)} \sim \mathcal{N}(0, \sigma_w^2) \quad (7)$$

$$b_i^{(l)} \sim \mathcal{N}(0, \sigma_b^2) \quad (8)$$

$$\sigma_y^2 \sim \text{InverseGamma}(\alpha, \beta) \quad (9)$$

Prior Considerations

Weight Scale: Controls capacity

- Small σ_w : Smooth functions
- Large σ_w : Complex functions

Architecture Prior:

- Number of layers
- Hidden units per layer
- Activation functions

Practical Tip

Use empirical Bayes or cross-validation for hyperparameter selection

Inference Methods for BNNs

1. Hamiltonian Monte Carlo

- **Exact sampling** (asymptotically)
- Uses gradient information
- Handles correlations well
- Computationally intensive

2. Variational Inference

- Approximate posterior $q(\mathbf{w}; \phi)$
- Minimize KL divergence
- Scalable to large networks
- Mean-field assumption

3. Monte Carlo Dropout

- Approximate Bayesian inference
- Keep dropout at test time
- Ensemble of sub-networks
- Computationally efficient

4. Ensemble Methods

- Train multiple networks
- Different initializations
- Bootstrap sampling
- Deep ensembles

VI Objective

$$\mathcal{L}(\phi) = \mathbb{E}_{q(\mathbf{w}; \phi)} [\log p(\mathcal{D} | \mathbf{w})] - \text{KL}(q(\mathbf{w}) || p(\mathbf{w}))$$

Practical Recommendation

Small networks: MCMC

Large networks: Variational inference

Variational Inference for BNNs: Bayes by Backprop

Algorithm: Bayes by Backprop (Blundell et al., 2015)

Algorithm 1 Variational Inference for BNN

- 1: **Initialize:** Variational parameters ϕ for $q(\mathbf{w}|\phi)$
 - 2: **for** each iteration **do**
 - 3: Sample weights: $\mathbf{w} \sim q(\mathbf{w}|\phi)$
 - 4: Compute loss: $\mathcal{L} = -\log p(\mathcal{D}|\mathbf{w}) + \text{KL}(q(\mathbf{w})\|p(\mathbf{w}))$
 - 5: Compute gradients: $\nabla_{\phi}\mathcal{L}$ using reparameterization trick
 - 6: Update: $\phi \leftarrow \phi - \alpha \nabla_{\phi}\mathcal{L}$
 - 7: **end for**
-

Reparameterization Trick:

$$\mathbf{w} = \boldsymbol{\mu} + \boldsymbol{\sigma} \odot \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(0, I) \quad (10)$$

$$q(\mathbf{w}) = \mathcal{N}(\mathbf{w}; \boldsymbol{\mu}, \text{diag}(\boldsymbol{\sigma}^2)) \quad (11)$$

Key Innovation: Gradient-based optimization of approximate posterior

BNN Applications and Case Studies

1. Regression with Uncertainty

- Heteroscedastic noise modeling
- Confidence intervals for predictions
- Outlier detection
- Active learning applications

2. Classification with Calibration

- Well-calibrated probabilities
- Uncertainty in predictions
- Out-of-distribution detection
- Medical diagnosis applications

3. Reinforcement Learning

- Exploration via uncertainty
- Thompson sampling

Case Study: Medical Diagnosis

Problem: Skin cancer classification

Dataset: 10,000 dermoscopy images

BNN Results:

- 94.2% accuracy (vs 93.8% standard NN)
- **Well-calibrated** confidence scores
- Identifies uncertain cases for expert review
- 15% reduction in misdiagnosis risk

Key Insight: Uncertainty quantification more valuable than accuracy gain

Industrial Impact

BNNs enable **safe AI deployment** in critical applications

Introduction to Gaussian Processes

Definition (Gaussian Process)

A Gaussian process is a collection of random variables, any finite number of which have a joint Gaussian distribution.

Key Insight: Instead of parameterizing functions, put distributions directly over functions.

$$f(\mathbf{x}) \sim \mathcal{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}')) \quad (12)$$

$$m(\mathbf{x}) = \mathbb{E}[f(\mathbf{x})] \quad (13)$$

$$k(\mathbf{x}, \mathbf{x}') = \text{Cov}[f(\mathbf{x}), f(\mathbf{x}')] \quad (14)$$

Properties:

- Non-parametric method
- Infinite dimensional

Finite Dimensional Consistency

For any finite set $\{\mathbf{x}_1, \dots, \mathbf{x}_n\}$:

$$\begin{pmatrix} f(\mathbf{x}_1) \\ \vdots \\ f(\mathbf{x}_n) \end{pmatrix} \sim \mathcal{N} \left(\begin{pmatrix} m(\mathbf{x}_1) \\ \vdots \\ m(\mathbf{x}_n) \end{pmatrix}, K \right)$$

where $K_{ij} = k(\mathbf{x}_i, \mathbf{x}_j)$

Computational Complexity

Kernel Functions and Prior Specification

Popular Kernel Functions:

1. Squared Exponential (RBF):

$$k(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \exp \left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\ell^2} \right)$$

2. Matérn:

$$k(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \frac{2^{1-\nu}}{\Gamma(\nu)} \left(\frac{\sqrt{2\nu}r}{\ell} \right)^\nu K_\nu \left(\frac{\sqrt{2\nu}r}{\ell} \right)$$

3. Periodic:

$$k(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \exp \left(-\frac{2 \sin^2(\pi|x - x'|/p)}{\ell^2} \right)$$

4. Linear:

Kernel Properties:

- **Length scale** ℓ : Controls smoothness
- **Signal variance** σ_f^2 : Output scale
- **Noise variance** σ_n^2 : Observation noise

Kernel Composition:

- **Addition:** $k_1 + k_2$ (combining patterns)
- **Multiplication:** $k_1 \times k_2$ (conjunction)
- **Scaling:** αk (amplitude)

Example: Trend + Periodic

$$k(\mathbf{x}, \mathbf{x}') = k_{\text{linear}}(\mathbf{x}, \mathbf{x}') + k_{\text{periodic}}(\mathbf{x}, \mathbf{x}') + k_{\text{noise}}(\mathbf{x}, \mathbf{x}')$$

GP Regression: Predictive Distribution

Training Data: $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$

Likelihood: $y_i = f(\mathbf{x}_i) + \epsilon_i$, where $\epsilon_i \sim \mathcal{N}(0, \sigma_n^2)$

Predictive Distribution

For a new input \mathbf{x}^* , the predictive distribution is:

$$p(f^*|\mathbf{x}^*, \mathcal{D}) = \mathcal{N}(f^*; \mu^*, (\sigma^*)^2) \quad (15)$$

$$\mu^* = \mathbf{k}^T (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y} \quad (16)$$

$$(\sigma^*)^2 = k(\mathbf{x}^*, \mathbf{x}^*) - \mathbf{k}^T (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{k} \quad (17)$$

where:

- $\mathbf{k} = [k(\mathbf{x}^*, \mathbf{x}_1), \dots, k(\mathbf{x}^*, \mathbf{x}_n)]^T$
- $\mathbf{K}_{ij} = k(\mathbf{x}_i, \mathbf{x}_j)$
- $\mathbf{y} = [y_1, \dots, y_n]^T$

Key Properties:

- Exact Bayesian inference (given kernel)

Hyperparameter Learning in GPs

Marginal Likelihood:

$$p(\mathbf{y}|\mathbf{X}, \boldsymbol{\theta}) = \mathcal{N}(\mathbf{y}; \mathbf{0}, \mathbf{K} + \sigma_n^2 \mathbf{I}) \quad (18)$$

$$\log p(\mathbf{y}|\mathbf{X}, \boldsymbol{\theta}) = -\frac{1}{2} \mathbf{y}^T \mathbf{K}_y^{-1} \mathbf{y} \quad (19)$$

$$-\frac{1}{2} \log |\mathbf{K}_y| - \frac{n}{2} \log 2\pi \quad (20)$$

where $\mathbf{K}_y = \mathbf{K} + \sigma_n^2 \mathbf{I}$ and $\boldsymbol{\theta}$ are hyperparameters.

Optimization:

$$\boldsymbol{\theta}^* = \arg \max_{\boldsymbol{\theta}} \log p(\mathbf{y}|\mathbf{X}, \boldsymbol{\theta})$$

Three Terms Interpretation:

- ① **Data fit:** $-\frac{1}{2} \mathbf{y}^T \mathbf{K}_y^{-1} \mathbf{y}$
- ② **Complexity penalty:** $-\frac{1}{2} \log |\mathbf{K}_y|$
- ③ **Normalization:** $-\frac{n}{2} \log 2\pi$

Automatic Occam's Razor

GPs automatically balance model complexity and data fit through the marginal likelihood

Practical Implementation

Use gradient-based optimization (L-BFGS) with multiple random restarts

Sparse Gaussian Processes

Problem: Standard GP inference scales as $O(n^3)$

Solution: Inducing point methods

Key Idea: Approximate full GP using $m \ll n$ inducing points

$$f(\mathbf{x}) \approx \tilde{f}(\mathbf{x}) = \mathbf{k}_*^T \mathbf{K}_{mm}^{-1} \mathbf{f}_m \quad (21)$$

$$\mathbf{f}_m = [f(\mathbf{z}_1), \dots, f(\mathbf{z}_m)]^T \quad (22)$$

where $\{\mathbf{z}_i\}_{i=1}^m$ are inducing inputs.

Variational Sparse GP:

- Optimize inducing inputs $\{\mathbf{z}_i\}$
- Variational distribution $q(\mathbf{f}_m)$

FITC Approximation

Fully Independent Training Conditional:

$$q(\mathbf{f}) = p(\mathbf{f}|\mathbf{f}_m)q(\mathbf{f}_m)$$

Stochastic Variational GP

- Mini-batch training
- Natural gradients
- Scales to millions of points
- $O(m^3)$ per iteration

Modern Extensions:

1. Bayesian Optimization

- Expensive function optimization
- Acquisition functions (EI, UCB, PI)
- Hyperparameter tuning
- Experimental design

2. Time Series Forecasting

- Temporal kernels
- Uncertainty in predictions
- Missing data handling
- Irregular time series

3. Spatial Statistics

- Geostatistics and kriging
- Environmental monitoring

Case Study: Drug Discovery

Problem: Optimize molecular properties

Setup:

- 10,000 molecules tested
- Each test costs \$1000
- Goal: Find top 1% molecules

GP-based Optimization:

- Molecular fingerprints as features
- Tanimoto kernel for similarity
- Expected improvement acquisition

Results:

- 95% reduction in tests needed
- Found optimal molecules in 500 tests

Deep Gaussian Processes

Motivation: Combine flexibility of deep learning with uncertainty of GPs

Inference Challenges:

- Intractable posterior
- Doubly stochastic variational inference
- Reparameterization trick for GPs
- Computational complexity

Architecture:

$$\mathbf{h}_1 \sim \mathcal{GP}(\mathbf{0}, k_1(\mathbf{x}, \mathbf{x}')) \quad (23)$$

$$\mathbf{h}_2 \sim \mathcal{GP}(\mathbf{0}, k_2(\mathbf{h}_1, \mathbf{h}'_1)) \quad (24)$$

$$\vdots \quad (25)$$

$$\mathbf{y} \sim \mathcal{GP}(\mathbf{0}, k_L(\mathbf{h}_{L-1}, \mathbf{h}'_{L-1})) \quad (26)$$

Properties:

- Non-stationary kernels
- Hierarchical feature learning
- Uncertainty propagation through layers
- Automatic relevance determination

Variational Approach

- Variational distribution for each layer
- Monte Carlo estimates
- Natural gradient optimization

Applications:

- High-dimensional regression

Neural Networks vs Gaussian Processes: The Infinite Width Limit

Remarkable Connection: Neural networks converge to Gaussian processes in the infinite width limit

Theorem (Neal, 1996)

Consider a single hidden layer neural network:

$$f(\mathbf{x}) = \frac{1}{\sqrt{H}} \sum_{i=1}^H v_i \sigma(w_i^T \mathbf{x} + b_i)$$

As $H \rightarrow \infty$ with i.i.d. weights w_i, v_i, b_i , the function $f(\mathbf{x})$ converges to a Gaussian process.

Neural Tangent Kernel:

- Describes infinite-width NN dynamics
- Fixed kernel during training

Modern Extensions

Deep Neural Networks:

- Infinite depth limit
- Bayesian neural networks at scale

Variational Inference: Connecting BNNs and GPs

Unified Framework: Both BNNs and GPs can be viewed through variational inference lens
Modern Developments:

Bayesian Neural Networks:

$$\text{ELBO} = \mathbb{E}_{q(\mathbf{w})}[\log p(\mathcal{D}|\mathbf{w})] - \text{KL}(q(\mathbf{w})\|p(\mathbf{w})) \quad (27)$$

Sparse Gaussian Processes:

$$\text{ELBO} = \mathbb{E}_{q(\mathbf{f})}[\log p(\mathcal{D}|\mathbf{f})] - \text{KL}(q(\mathbf{f})\|p(\mathbf{f})) \quad (28)$$

Common Techniques:

- Reparameterization trick
- Natural gradients
- Stochastic optimization

Normalizing Flows

- Flexible posterior approximations
- Invertible transformations
- Better than mean-field

Neural Processes

- Combine NNs and GPs
- Amortized inference
- Meta-learning for functions

Practical Tools:

- PyTorch, TensorFlow Probability

Software Ecosystem for Bayesian ML

Framework	Strengths	Applications	Language
PyTorch	Flexible, research-friendly	Custom BNN architectures	Python
TensorFlow Prob.	Production-ready, scalable	Large-scale deployment	Python
GPyTorch	GPU acceleration for GPs	Large-scale GP inference	Python
GPflow	TensorFlow-based GPs	Deep GPs, sparse methods	Python
Stan	Probabilistic programming	Custom model specification	Multiple
PyMC	User-friendly Bayesian	Educational, prototyping	Python

Performance Considerations:

Production Deployment:



Implementation Example: Simple BNN in PyTorch

Bayesian Linear Layer

```
class BayesianLinear(nn.Module):
    def __init__(self, in_features, out_features):
        super().__init__()
        # Weight parameters
        self.weight_mu = nn.Parameter(
            torch.zeros(out_features, in_features))
        self.weight_rho = nn.Parameter(
            torch.ones(out_features, in_features) * -3)

        # Bias parameters
        self.bias_mu = nn.Parameter(
            torch.zeros(out_features))
        self.bias_rho = nn.Parameter(
            torch.ones(out_features) * -3)
```

Key Components:

- Variational parameters: μ, ρ
- Reparameterization: $w = \mu + \sigma \epsilon$
- Softplus: $\sigma = \log(1 + \exp(\rho))$

Training Loop:

- Sample weights each forward pass
- Compute ELBO loss
- Backprop through sampling
- Update variational parameters

Best Practices and Common Pitfalls

Bayesian Neural Networks:

Best Practices

- Start with prior sensitivity analysis
- Use multiple random seeds
- Monitor KL divergence during training
- Validate uncertainty calibration
- Consider computational budget

Common Pitfalls

- Overconfident posterior approximations
- Poor initialization of variational parameters
- Ignoring computational overhead

Gaussian Processes:

Best Practices

- Choose kernels based on problem structure
- Use multiple optimization restarts
- Validate on held-out data
- Consider sparse approximations for large data
- Monitor numerical stability

Common Pitfalls

- Poor kernel choice for the problem
- Local optima in hyperparameter optimization
- Numerical issues with matrix inversion

Current Research Frontiers

Methodological Advances:

- **Continual learning** with uncertainty
- **Federated Bayesian learning**
- **Physics-informed priors**
- **Causal discovery** with GPs
- **Multi-modal** Bayesian models

Computational Innovations:

- **Quantum-enhanced** sampling
- **Neuromorphic computing** for BNNs
- **Distributed inference** at scale
- **Edge deployment** of Bayesian models

Application Domains:

- **Autonomous systems** (vehicles, drones)
- **Healthcare AI** with safety guarantees
- **Climate modeling** and prediction
- **Financial risk** assessment
- **Scientific discovery** acceleration

Societal Impact:

- **Trustworthy AI** development
- **Algorithmic fairness** with uncertainty
- **Privacy-preserving** ML
- **Explainable AI** through Bayesian lens

The Next Decade

Bayesian ML will become the **standard approach** for safety-critical applications requiring

Summary and Key Takeaways

Bayesian Neural Networks:

- **Uncertainty quantification** for deep learning
- **Automatic regularization** through priors
- **Computational challenges** in large networks
- **Growing adoption** in critical applications

Gaussian Processes:

- **Non-parametric** Bayesian approach
- **Exact inference** for regression
- **Flexible** through kernel design
- **Scalability** remains a challenge

Practical Guidelines:

- Start with **simple baselines**
- Validate **uncertainty calibration**
- Consider **computational constraints**
- Use **appropriate software** tools
- Focus on **problem-specific** solutions

When to Use What?

BNNs: Large datasets, complex patterns, representation learning

GPs: Small-medium datasets, interpretability, principled uncertainty

Thank You

Questions & Discussion

Diogo Ribeiro

ESMAD – Escola Superior de Média Arte e Design

Lead Data Scientist, Mysense.ai

dfr@esmad.ipp.pt

<https://orcid.org/0009-0001-2022-7072>

Slides and code available at:

github.com/diogoribeiro7/academic-presentations