

# Feature Engineering & Selection

## From Raw Data to ML-Ready Features

Diogo Ribeiro  
ESMAD – Escola Superior de Média Arte e Design  
Lead Data Scientist, Mysense.ai

October 29, 2025

# Outline

- 1 Introduction and Philosophy
- 2 Data Types and Transformations
- 3 Advanced Feature Creation
- 4 Dimensionality Reduction
- 5 Feature Selection Methods
- 6 Practical Considerations and Best Practices
- 7 Summary and Next Steps

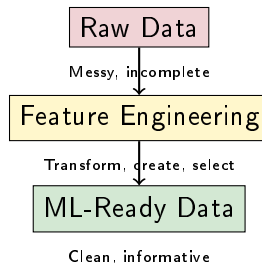
# The Art and Science of Feature Engineering

*"Coming up with features is difficult, time-consuming, requires expert knowledge. 'Applied machine learning' is basically feature engineering."*

– Andrew Ng

## Why Features Matter:

- **Garbage in, garbage out:** Poor features  $\Rightarrow$  poor models
- **Domain knowledge:** Good features encode expert insights
- **Model performance:** Often more impact than algorithm choice
- **Interpretability:** Good features are meaningful to humans



**80% of  
DS time!**

## The 80-20 Rule

80% of data science time is spent on data preparation and feature engineering, 20% on modeling.

# Feature Engineering Pipeline Overview

## The Complete Pipeline:

### 1 Data Understanding

- Exploratory data analysis
- Data quality assessment
- Domain knowledge integration

### 2 Cleaning & Preprocessing

- Missing value handling
- Outlier detection/treatment
- Data type conversions

### 3 Feature Creation

- Transformations
- Interactions
- Domain-specific features

### 4 Feature Selection

- Remove redundant features
- Statistical significance

## Success Metrics

- **Model Performance:** Accuracy, AUC, RMSE
- **Interpretability:** Can humans understand features?
- **Stability:** Robust to new data
- **Efficiency:** Fast to compute and store

## Common Mistakes

- **Data leakage:** Using future information
- **Overfitting:** Too many features for sample size
- **Domain ignorance:** Features that don't make sense

# Understanding Your Data Types

Data Type	Examples	Challenges	Common Transformations
Numerical	Age, income, temperature	Skewness, outliers, scale	Log, square root, standardization
Categorical	Color, country, brand	High cardinality, ordering	One-hot, label encoding, embeddings
Ordinal	Education level, ratings	Preserving order	Ordinal encoding, polynomial features
Temporal	Timestamps, dates	Seasonality, trends	Date parts, lags, rolling statistics
Text	Reviews, descriptions	High dimensionality	TF-IDF, embeddings, sentiment
Geospatial	Coordinates, addresses	Projection, distance	Distance features, clustering

# Numerical Feature Transformations

## Common Issues with Numerical Features:

```
import numpy as np
import pandas as pd
from scipy import stats
from sklearn.preprocessing import (
    StandardScaler, MinMaxScaler,
    RobustScaler, PowerTransformer,
    QuantileTransformer
)

# Example: Highly skewed income data
np.random.seed(42)
income = np.random.lognormal(10, 1, 1000)
print(f"Skewness: {stats.skew(income):.2f}")
print(f"Range: [{income.min():.0f}, {income.
    max():.0f}]")

# Transformation strategies
transformations = [
```

## Scaling Strategies:

- **StandardScaler:**  $z = \frac{x - \mu}{\sigma}$ 
  - Good for normally distributed data
  - Sensitive to outliers
- **MinMaxScaler:**  $x' = \frac{x - \min(x)}{\max(x) - \min(x)}$ 
  - Bounds data to [0,1]
  - Very sensitive to outliers
- **RobustScaler:** Uses median and IQR

# Categorical Feature Encoding

**The Categorical Challenge:** ML algorithms need numbers, not categories.

```
import pandas as pd
from sklearn.preprocessing import (
    LabelEncoder, OneHotEncoder,
    OrdinalEncoder
)
from category_encoders import (
    TargetEncoder, BinaryEncoder,
    HashingEncoder, LeaveOneOutEncoder
)

# Sample categorical data
data = pd.DataFrame({
    'color': ['red', 'blue', 'green', 'red', 'blue'],
    'size': ['small', 'medium', 'large', 'medium', 'small'],
    'brand': ['nike', 'adidas', 'puma', 'nike', '']
```

## Encoding Strategy Guide:

Method	Cardinality	Best For
Label	Any	Ordinal data
One-Hot	Low (< 10)	Nominal data
Target	Medium	High predictive power

# Temporal Feature Engineering

**Time Series Features:** Extract meaningful patterns from timestamps.

```
import pandas as pd
import numpy as np
from datetime import datetime, timedelta

# Sample time series data
dates = pd.date_range(start='2020-01-01', end='
    2023-12-31', freq='D')
np.random.seed(42)
sales = 100 + 10*np.sin(2*np.pi*np.arange(len(
    dates))/365.25) + \
    5*np.random.randn(len(dates))

df = pd.DataFrame({'date': dates, 'sales': sales})
df['date'] = pd.to_datetime(df['date'])

# Basic temporal features
df['year'] = df['date'].dt.year
```

## Temporal Feature Categories:

- **Date Parts:** Year, month, day, hour
- **Cyclical:** Sin/cos encoding for



# Polynomial and Interaction Features

## Capturing Non-linear Relationships:

```
import numpy as np
import pandas as pd
from sklearn.preprocessing import
    PolynomialFeatures
from sklearn.model_selection import
    cross_val_score
from sklearn.linear_model import
    LinearRegression
from sklearn.pipeline import Pipeline

# Generate sample data with interactions
np.random.seed(42)
n_samples = 1000
X1 = np.random.randn(n_samples)
X2 = np.random.randn(n_samples)
X3 = np.random.randn(n_samples)
```

## When to Use Polynomial Features:

- **Linear models:** Add non-linearity
- **Known relationships:** Domain knowledge

# Domain-Specific Feature Engineering

## Case Study: E-commerce Customer Features

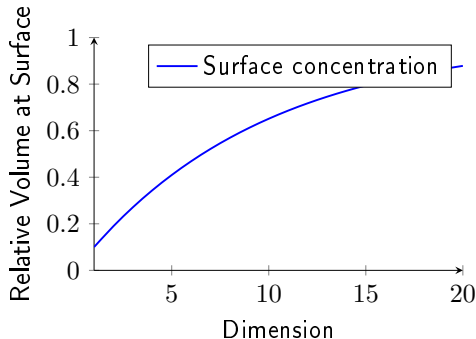
```
import pandas as pd
import numpy as np
from datetime import datetime, timedelta

# Sample e-commerce transaction data
np.random.seed(42)
transactions = pd.DataFrame({
    'customer_id': np.repeat(range(1000), 10),
    'transaction_date': pd.to_datetime('2023-01-01') +
                        pd.to_timedelta(np.random.
randint(0, 365, 10000), unit='D'),
    'amount': np.random.exponential(50, 10000),
    'product_category': np.random.choice(['electronics', 'clothing', 'books', 'home'],
10000),
    'is_returned': np.random.binomial(1, 0.1,
```

# The Curse of Dimensionality

## Why High Dimensions Are Problematic:

- **Sparsity:** Data points become sparse in high-D space
- **Distance concentration:** All points equidistant
- **Overfitting:** More parameters than samples
- **Computational cost:**  $O(d^k)$  complexity
- **Visualization:** Impossible to plot  $> 3D$



**Mathematical Intuition:** Volume of unit hypersphere in  $d$  dimensions:

$$V_d = \frac{\pi^{d/2}}{\Gamma(d/2 + 1)}$$

## The $d \gg n$ Problem

When dimensions exceed samples:

- Perfect separation possible

# Principal Component Analysis (PCA)

**Goal:** Find linear combinations of features that maximize variance.

## Mathematical Formulation:

Maximize:  $\text{Var}(\mathbf{w}^T \mathbf{X})$

Subject to:  $\|\mathbf{w}\|^2 = 1$

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.decomposition import PCA
4 from sklearn.datasets import
    load_breast_cancer
5 from sklearn.preprocessing import
    StandardScaler
6
7 # Load high-dimensional dataset
8 data = load_breast_cancer()
9 X, y = data.data, data.target
10
11 print(f"Original dimensions: {X.shape}")
12
13 # Standardize features (crucial for PCA)
14 scaler = StandardScaler()
15 X_scaled = scaler.fit_transform(X)
```

# Advanced Dimensionality Reduction Techniques

## Non-linear Methods for Complex Data:

```
import numpy as np
from sklearn.manifold import TSNE
from umap import UMAP
from sklearn.decomposition import KernelPCA
from sklearn.datasets import make_swiss_roll

# Generate non-linear data
X, color = make_swiss_roll(n_samples=1000, noise
                           =0.1)

# 1. Kernel PCA (non-linear PCA)
kpca_rbf = KernelPCA(n_components=2, kernel='rbf',
                     gamma=0.1)
X_kpca = kpca_rbf.fit_transform(X)


# 2. t-SNE (preserves local structure)
tsne = TSNE(n_components=2, perplexity=30,
```

Method Comparison:

# Feature Selection: Why and When

## Why Remove Features?

- **Overfitting**: Too many features relative to samples
- **Noise**: Irrelevant features hurt performance
- **Multicollinearity**: Redundant information
- **Computational cost**: Storage and processing
- **Interpretability**: Simpler models are easier to understand



## The Bias-Variance Perspective:

Fewer features  $\Rightarrow$  Higher bias, Lower variance  
(5)

## When to Apply Feature Selection

- High-dimensional data ( $p \gg n$ )
- Noisy or redundant features
- Interpretability requirements

# Filter Methods: Statistical Feature Selection

**Idea:** Rank features by statistical measures, independent of the learning algorithm.

```
import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.feature_selection import (
    SelectKBest, f_classif, chi2,
    mutual_info_classif,
    VarianceThreshold,
    SelectPercentile
)
from sklearn.preprocessing import StandardScaler
from scipy.stats import pearsonr

# Generate sample data with irrelevant features
X, y = make_classification(
    n_samples=1000, n_features=20,
    n_informative=5, n_redundant=3,
    n_clusters_per_class=1, random_state=42
)
```

# Wrapper Methods: Model-Based Selection

**Idea:** Use the target model's performance to evaluate feature subsets.

```
import numpy as np
from sklearn.model_selection import
    cross_val_score
from sklearn.ensemble import
    RandomForestClassifier
from sklearn.feature_selection import (
    RFE, RFECV, SequentialFeatureSelector
)
from sklearn.linear_model import
    LogisticRegression
from sklearn.datasets import make_classification

# Generate data
X, y = make_classification(
    n_samples=500, n_features=15,
    n_informative=5, n_redundant=3,
    random_state=42
```



# Embedded Methods: Built-in Feature Selection

**Idea:** Feature selection is integrated into the model training process.

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_regression
from sklearn.linear_model import Lasso, Ridge,
    ElasticNet
from sklearn.ensemble import RandomForestRegressor
from sklearn.feature_selection import
    SelectFromModel
from sklearn.model_selection import
    cross_val_score

# Generate regression data with some irrelevant
# features
X, y = make_regression(
    n_samples=200, n_features=20,
    n_informative=5, noise=10,
    random_state=42
```

# Data Leakage: The Silent Killer

**What is Data Leakage?** Information from the future or target variable inadvertently included in features.

## Types of Leakage:

- **Temporal leakage:** Using future information
- **Target leakage:** Features that contain target info
- **Train-test leakage:** Data preprocessing on full dataset

## Common Examples:

- Credit scoring using payment history after loan decision
- Medical diagnosis using treatment

## Preventing Leakage:

- 1 **Time-aware splits:**
  - Train on past, test on future
  - No random shuffling for time series
- 2 **Proper preprocessing:**
  - Fit transformers only on training data
  - Apply to test data, don't refit
- 3 **Domain knowledge:**
  - What information is available when?
  - Business process understanding
- 4 **Feature audit:**
  - Check correlations with target
  - Validate with domain experts

**Golden Rule**

# Pipeline Design and Reproducibility

**Best Practice:** Use sklearn pipelines for reproducible feature engineering.

```
import pandas as pd
import numpy as np
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import (
    StandardScaler, OneHotEncoder,
    FunctionTransformer
)
from sklearn.feature_selection import SelectKBest,
    f_regression
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import cross_val_score

# Sample mixed-type dataset
np.random.seed(42)
data = pd.DataFrame({
    'numeric1': np.random.randn(1000),
```

# Feature Engineering Checklist

## Data Understanding

- ☐ Exploratory data analysis
- ☐ Missing value patterns
- ☐ Outlier detection
- ☐ Data type validation
- ☐ Domain expert consultation

## Feature Creation

- ☐ Handle missing values appropriately
- ☐ Scale/normalize numerical features
- ☐ Encode categorical variables
- ☐ Create interaction features
- ☐ Extract temporal features
- ☐ Engineer domain-specific features

## Leakage Prevention

- ☐ Time-aware data splitting
- ☐ Fit transformers only on training data
- ☐ Audit features for target information
- ☐ Validate business logic
- ☐ Check correlation with target

## Validation & Testing

- ☐ Cross-validation with proper splits
- ☐ Test on unseen data
- ☐ Monitor feature distributions
- ☐ A/B test in production
- ☐ Track model performance over time

# Key Takeaways

## Core Principles:

- **Domain knowledge** is as important as technical skills
- **Systematic approach** beats ad-hoc feature creation
- **Validation** is crucial for avoiding overfitting
- **Pipelines** ensure reproducibility and prevent leakage

## Technical Skills Learned:

- Data type-specific transformations
- Categorical encoding strategies
- Dimensionality reduction techniques

## Common Pitfalls Avoided:

- **Data leakage** through improper preprocessing
- **Overfitting** with too many features
- **Scale issues** in mixed-type data
- **Target leakage** in feature creation
- **Irreproducible** results

## The Art vs Science

**Science:** Statistical methods, validation procedures, systematic evaluation

**Art:** Domain insights, creative feature combinations, business intuition

# Next Steps in Your Data Science Journey

## Immediate Next Topics:

### ① Causal Inference for Data Scientists

- Moving beyond correlation
- Experimental design
- Observational causal methods

### ② Explainable AI & Model Interpretability

- SHAP and LIME
- Global vs local explanations
- Building trust in models

### ③ Experimental Design & A/B Testing

## Practice Projects:

- Build end-to-end feature engineering pipeline
- Kaggle competition with focus on feature engineering
- Industry-specific feature creation
- Automated feature engineering tools

## Advanced Topics to Explore:

- Automated feature engineering (Featuretools)
- Deep feature synthesis
- Graph-based features
- Text feature engineering

# Thank You

## Questions & Discussion

**Diogo Ribeiro**

ESMAD – Escola Superior de Média Arte e Design  
Lead Data Scientist, Mysense.ai

dfr@esmad.ipp.pt

<https://orcid.org/0009-0001-2022-7072>

*Slides and code available at:*

[github.com/diogoribeiro7/academic-presentations](https://github.com/diogoribeiro7/academic-presentations)

*Next: Causal Inference for Data Scientists*