

# Markov Chain Monte Carlo

## Theory, Modern Algorithms, and Applications

Diogo Ribeiro  
ESMAD – Escola Superior de Média Arte e Design  
Lead Data Scientist, Mysense.ai

October 27, 2025

# Outline

- 1 Introduction and Foundations
- 2 Mathematical Foundations
- 3 Classical Algorithms
- 4 Modern Algorithms
- 5 Diagnostics and Convergence
- 6 Applications
- 7 Implementation and Software
- 8 Future Directions

# The Monte Carlo Revolution

## Historical Timeline:

- **1940s:** Stanisław Ulam – Manhattan Project
- **1953:** Metropolis et al. – First MCMC algorithm
- **1970:** Hastings – Generalized acceptance criterion
- **1984:** Geman & Geman – Gibbs sampling
- **1987:** Duane et al. – Hamiltonian Monte Carlo
- **2011:** Hoffman & Gelman – NUTS algorithm

## Revolutionary Impact

- Transformed statistics from analytical to computational
- Enabled Bayesian inference for complex models
- Made high-dimensional problems tractable
- Foundation of modern ML/AI

## Key Insight

MCMC didn't just change *how* we compute – it changed *what* we can compute.

# The Fundamental Sampling Challenge

**The Problem:** Sample from  $\pi(\mathbf{x})$  when direct methods fail

## Traditional Methods Breakdown:

- **Inverse transform:** No closed form CDF
- **Rejection sampling:** Exponential inefficiency in high-D
- **Grid methods:** Curse of dimensionality
- **Importance sampling:** Poor proposal overlap

## MCMC Solution

Construct Markov chain with  $\pi(\mathbf{x})$  as stationary distribution

Method	Dimension	Feasible?
Direct sampling	1D	Yes
Inverse transform	Simple PDFs	Yes
Rejection sampling	> 5D	No
Grid methods	> 3D	No
<b>MCMC</b>	<b>Any</b>	<b>Yes</b>

## Critical Applications:

- Bayesian neural networks ( $10^6+$  parameters)
- Financial risk models
- Climate simulations
- Phylogenetic inference

# Markov Chain Theory

## Definition (Markov Chain)

A sequence  $\{X_n\}_{n \geq 0}$  is a Markov chain if:

$$P(X_{n+1} = x_{n+1} | X_0, \dots, X_n) = P(X_{n+1} = x_{n+1} | X_n)$$

## Key Concepts:

- **Transition kernel:**

$$P(x, A) = P(X_{n+1} \in A | X_n = x)$$

- **Chapman-Kolmogorov:**

$$P^n(x, A) = \int P^{n-1}(x, dy) P(y, A)$$

- **Invariant distribution:**

$$\pi(A) = \int \pi(dx) P(x, A)$$

## Theorem (Ergodic Theorem)

*If the chain is irreducible, aperiodic, and positive recurrent, then:*

$$\lim_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n f(X_i) = \int f(x) \pi(dx) \quad a.s.$$

## Detailed Balance Condition

$$\pi(x) P(x, dy) = \pi(y) P(y, dx)$$

# Convergence Theory and Rates

## Theorem (Convergence to Stationarity)

*Under regularity conditions:*

$$\|P^n(x, \cdot) - \pi(\cdot)\|_{TV} \leq C\rho^n$$

*where  $\rho < 1$  is the second-largest eigenvalue.*

### Mixing Time:

$$\tau_{mix}(\epsilon) = \min\{n : \max_x \|P^n(x, \cdot) - \pi(\cdot)\|_{TV} \leq \epsilon\}$$

**Spectral Gap:**  $\gamma = 1 - \rho$  determines convergence rate

## Factors Affecting Convergence

- **Geometry:** Condition number of target
- **Dimensionality:** Concentration phenomena
- **Multimodality:** Barrier crossing
- **Step size:** Acceptance vs exploration trade-off

## Central Limit Theorem for MCMC

$$\sqrt{n}(\bar{f}_n - \pi(f)) \xrightarrow{d} N(0, \sigma_f^2)$$

where  $\sigma_f^2 = \text{Var}_\pi(f) + 2 \sum_{k=1}^{\infty} \text{Cov}_\pi(f(X_0), f(X_k))$

# Metropolis-Hastings Algorithm

---

## Algorithm Metropolis-Hastings

---

- 1: Initialize  $x^{(0)}$
- 2: **for**  $t = 0, 1, 2, \dots$  **do**
- 3:   Propose  $y \sim q(y|x^{(t)})$
- 4:   Compute acceptance probability:

$$\alpha(x^{(t)}, y) = \min \left( 1, \frac{\pi(y)q(x^{(t)}|y)}{\pi(x^{(t)})q(y|x^{(t)})} \right)$$

- 5:   Accept  $x^{(t+1)} = y$  with probability  $\alpha$ , otherwise  $x^{(t+1)} = x^{(t)}$
  - 6: **end for**
- 

## Key Variants:

- **Random Walk:**  $q(y|x) = q(y - x)$

## Optimal Scaling Theory

For  $d$ -dimensional Gaussian targets:

## Adaptive Metropolis:

$$C_n = \frac{s_d}{n} \sum_{i=1}^n (X_i - \bar{X}_n)(X_i - \bar{X}_n)^T + s_d \epsilon_n I_d$$

where  $s_d = (2.38)^2/d$  and  $\epsilon_n \rightarrow 0$ .

## Advantages:

- Automatic tuning
- Adapts to target geometry
- Maintains ergodicity

## Multiple-try Metropolis:

- 1 Generate  $k$  proposals:  $y_1, \dots, y_k \sim q(\cdot|x)$
- 2 Select  $y_j$  with probability  $\propto \pi(y_j)$
- 3 Generate reference set from  $y_j$
- 4 Accept/reject based on ratio of weights

## Benefits

- Higher acceptance rates
- Better exploration
- Parallelizable proposals



# Gibbs Sampling and Blocking

## Standard Gibbs Sampling:

$$X_1^{(t+1)} \sim \pi(x_1 | X_2^{(t)}, X_3^{(t)}, \dots, X_d^{(t)}) \quad (1)$$

$$X_2^{(t+1)} \sim \pi(x_2 | X_1^{(t+1)}, X_3^{(t)}, \dots, X_d^{(t)}) \quad (2)$$

$$\vdots \quad (3)$$

$$X_d^{(t+1)} \sim \pi(x_d | X_1^{(t+1)}, X_2^{(t+1)}, \dots, X_{d-1}^{(t+1)}) \quad (4)$$

## Blocking Strategies:

- **Random scan:** Update components randomly
- **Block Gibbs:** Update correlated components together
- **Collapsed Gibbs:** Integrate out auxiliary variables

## Performance Considerations

- **Slow mixing:** High posterior correlations
- **Fast mixing:** Near-independence
- **Curse:**  $O(d^2)$  scaling with correlation

## Acceleration Techniques:

# Hamiltonian Monte Carlo

**Physical Intuition:** Frictionless particle on curved surface

**Hamiltonian System:**

$$H(q, p) = U(q) + K(p) \quad (5)$$

$$U(q) = -\log \pi(q) \text{ (potential energy)} \quad (6)$$

$$K(p) = \frac{1}{2} p^T M^{-1} p \text{ (kinetic energy)} \quad (7)$$

**Hamilton's equations:**

$$\frac{dq}{dt} = \frac{\partial H}{\partial p} = M^{-1} p \quad (8)$$

$$\frac{dp}{dt} = -\frac{\partial H}{\partial q} = -\nabla U(q) \quad (9)$$

**Leapfrog Integrator**

$$p_{t+\epsilon/2} = p_t - \frac{\epsilon}{2} \nabla U(q_t) \quad (10)$$

$$q_{t+\epsilon} = q_t + \epsilon M^{-1} p_{t+\epsilon/2} \quad (11)$$

$$p_{t+\epsilon} = p_{t+\epsilon/2} - \frac{\epsilon}{2} \nabla U(q_{t+\epsilon}) \quad (12)$$

# No-U-Turn Sampler (NUTS)

**Problem with HMC:** Manual tuning of step size  $\epsilon$  and number of steps  $L$

---

## Algorithm NUTS Algorithm (Simplified)

---

- 1: Sample momentum  $p_0 \sim N(0, M)$
  - 2: Set  $q_- = q_+ = q_0$ ,  $p_- = p_+ = p_0$
  - 3: Build trajectory by doubling until U-turn criterion:
  - 4: **while** no U-turn detected **do**
  - 5:   Double trajectory length in random direction
  - 6:   Check stopping criterion:  $(q_+ - q_-) \cdot p_+ < 0$  or  $(q_+ - q_-) \cdot p_- < 0$
  - 7: **end while**
  - 8: Sample uniformly from valid points in trajectory
- 

## Automatic Adaptation:

- **Step size:** Dual averaging to target acceptance rate

## Performance Benefits

- No manual parameter tuning

# Advanced MCMC Techniques

## Parallel Tempering:

- Multiple chains at different "temperatures"
- $\pi_i(x) \propto \pi(x)^{1/T_i}$  where  $T_1 < T_2 < \dots < T_k$
- Periodic swaps between chains
- Facilitates mode jumping

## Reversible Jump MCMC:

- Variable dimension problems
- Model selection applications
- Birth-death processes
- Careful design of dimension-changing moves

## Riemannian Manifold MCMC:

- Exploit geometric structure of parameter space
- Metric tensor:  $G(q) = \nabla^2 U(q)$  (Fisher information)
- Natural gradient directions
- Invariant to reparameterization

## Emerging Techniques

- **Neural MCMC:** Deep learning proposals
- **Quantum MCMC:** Quantum annealing
- **Piecewise deterministic:** Event-driven sampling

## Quantitative Diagnostics:

### 1. Gelman-Rubin Statistic:

$$\hat{R} = \sqrt{\frac{\hat{V}}{W}}$$

where  $\hat{V} = \frac{n-1}{n}W + \frac{1}{n}B$  and  $B, W$  are between/within chain variances.

**Target:**  $\hat{R} \leq 1.01$

### 2. Effective Sample Size:

$$ESS = \frac{mn}{1 + 2 \sum_{t=1}^T \hat{\rho}_t}$$

### 3. Monte Carlo Standard Error:

$$MCSE = \frac{\hat{\sigma}}{\sqrt{ESS}}$$

## Visual Diagnostics:

- **Trace plots:** Mixing and stationarity
- **Rank plots:** Chain uniformity
- **Autocorrelation:** Dependence structure
- **Energy plots:** HMC-specific diagnostics

## Diagnostic Protocol

- 1 Multiple chains, dispersed starts
- 2 Adequate warm-up ( $\geq 50\%$  samples)

# Common Convergence Problems

Problem	Symptoms	Solutions
Poor mixing	High autocorr., low ESS	Adaptive proposals, reparameterization, HMC
Multimodality	Chains in different modes	Parallel tempering, longer runs, multiple starts
Label switching	Erratic parameter traces	Post-processing, identifiability constraints
Heavy tails	Slow convergence	Robust proposals, tempering
High dimension	All diagnostics poor	Dimension reduction, hierarchical models

## Bayesian Neural Networks:

$$p(\mathbf{w}|\mathcal{D}) \propto p(\mathcal{D}|\mathbf{w})p(\mathbf{w})$$

### Challenges:

- $10^6 +$  parameters
- High correlations
- Complex posterior geometry
- Computational constraints

### MCMC Solutions:

- Stochastic gradient MCMC
- Subsampling techniques
- Variational-MCMC hybrids

## Gaussian Processes:

$$f(\mathbf{x}) \sim \mathcal{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}'))$$

### Hyperparameter Inference:

- Length scales, noise variance
- Kernel parameters
- Inducing point locations

### Case Study: Drug Discovery

- Molecular property prediction
- Uncertainty quantification critical
- MCMC for hyperparameter posteriors
- Enables active learning strategies

# Financial Risk Modeling

## Stochastic Volatility Models:

$$r_t = \mu + \sqrt{h_t} \epsilon_t \quad (13)$$

$$\log h_{t+1} = \alpha + \beta \log h_t + \sigma \eta_t \quad (14)$$

## MCMC for Parameter Estimation:

- Latent volatility states  $\{h_t\}$
- Model parameters  $(\alpha, \beta, \sigma, \mu)$
- Non-Gaussian state space model

## Portfolio Risk Assessment:

- Multivariate copula models
- Tail dependence estimation
- Value-at-Risk calculation

**Regulatory Applications:** Basel III capital requirements, stress testing, model validation

## Case Study: Credit Risk

**Problem:** Bank portfolio with 10,000 loans

**Model:** Hierarchical default probabilities

$$\text{logit}(p_{ij}) = \alpha_j + \beta^T x_{ij} \quad (15)$$

$$\alpha_j \sim N(\mu_\alpha, \sigma_\alpha^2) \quad (16)$$

## MCMC Results:

- 95% VaR: \$127M
- Expected Shortfall: \$89M
- Sector-specific risk factors
- Uncertainty intervals for risk metrics



## Climate Modeling:

- Earth system model calibration
- Parameter uncertainty quantification
- Ensemble generation for projections
- Millions of differential equations

## Example: Ocean Circulation

$$\frac{\partial \mathbf{u}}{\partial t} + (\mathbf{u} \cdot \nabla) \mathbf{u} = -\frac{1}{\rho} \nabla p + \nu \nabla^2 \mathbf{u} \quad (17)$$

**MCMC for:** Mixing coefficients, boundary conditions, forcing parameters

## Phylogenetic Inference:

- Tree topology uncertainty
- Branch length estimation
- Molecular clock models
- Ancestral sequence reconstruction

## Computational Challenges

- Discrete tree space
- Likelihood calculations  $O(n^2)$
- Proposal design for trees
- Parallel computation strategies

## Medical Applications:

- Personalized treatment design

# Modern Software Frameworks

Framework	Language	Strengths	Best For
Stan	C++/R/Python	HMC/NUTS, Speed	General purpose
PyMC	Python	User-friendly, AD	Research, education
JAX/NumPyro	Python	GPU, JIT compilation	Large scale
TensorFlow Prob.	Python	Deep learning integration	ML applications
JAGS	R/C++	Flexibility	Traditional models

## Performance Considerations:

- **Automatic differentiation:** Essential for HMC
- **GPU acceleration:** Massive parallelization

## Production Deployment:

- **Containerization:** Docker, Kubernetes
- **Monitoring:** Real-time convergence tracking

# Implementation Example: Adaptive Metropolis

```
class AdaptiveMetropolis:
    def __init__(self, target_log_prob, dim):
        self.target_log_prob = target_log_prob
        self.dim = dim
        self.cov = np.eye(dim) * 0.1
        self.mean = np.zeros(dim)

    def sample(self, n_samples, x0):
        samples = np.zeros((n_samples, self.dim))
        samples[0] = x0
        current_x = x0
        current_logp = self.target_log_prob(x0)

        accepted = 0
        for i in range(1, n_samples):
            # Propose
            proposal = np.random.multivariate_normal(
                current_x, self.cov)
            proposal_logp = self.target_log_prob(proposal)

            # Accept/reject
            if (np.log(np.random.random()) <
                proposal_logp - current_logp):
                current_x = proposal
                current_logp = proposal_logp
                accepted += 1

        samples[i] = current_x
```

## Key Features:

- Automatic covariance adaptation
- Robust numerical implementation
- Performance monitoring
- Configurable adaptation schedule

## Extensions:

- Parallel chains
- Online adaptation
- Constraint handling
- Warm-up phase management

## Production Considerations

- Memory-efficient updates

# Emerging Trends and Research Frontiers

## Neural-Enhanced MCMC:

- Deep learning proposal distributions
- Normalizing flows for reparameterization
- Neural ODEs for continuous dynamics
- Learned acceptance criteria

## Quantum Computing Integration:

- Quantum annealing for optimization
- Variational quantum algorithms
- Quantum-classical hybrid methods
- Exponential speedup potential

## Geometric Methods:

- Information geometry

## Large-Scale Applications:

- Federated learning with MCMC
- Privacy-preserving inference
- Distributed posterior computation
- Edge computing deployment

## Next Decade Challenges

- **Scale:** Billions of parameters
- **Speed:** Real-time inference
- **Robustness:** Model misspecification
- **Automation:** Minimal human intervention

## Interdisciplinary Impact:

• AI safety and robustness

# The Road Ahead

## Methodological Priorities:

- ① **Adaptive algorithms:** Self-tuning, robust to problem structure
- ② **Scalable architectures:** Distributed, GPU-accelerated computing
- ③ **Quality assurance:** Automatic validation, error detection
- ④ **User interfaces:** Accessible to non-experts
- ⑤ **Integration:** Seamless ML/AI ecosystem compatibility

## Application Domains:

- Scientific computing at exascale
- Real-time decision making

## Vision for 2030

### MCMC will be:

- Fully automated
- Hardware-optimized
- Ubiquitously deployed
- Theoretically grounded
- Practically transformative

## Call to Action

- Contribute to open-source tools
- Bridge theory and practice
- Foster interdisciplinary collaboration

# Summary and Key Takeaways

## Theoretical Foundations:

- Markov chain theory provides rigorous framework
- Convergence rates depend on spectral properties
- Detailed balance ensures correct stationary distribution

## Algorithmic Advances:

- HMC/NUTS: State-of-the-art for continuous distributions
- Adaptive methods: Automatic parameter tuning
- Specialized techniques: Problem-specific solutions

## Implementation Principles:

- Comprehensive diagnostics are essential
- Software frameworks enable productivity
- Performance optimization requires expertise

## The MCMC Paradigm

**Theory** → **Algorithms** → **Software** → **Applications**

## Future Outlook:

- Integration with AI/ML continues

# Thank You

**Diogo Ribeiro**

ESMAD – Escola Superior de Média Arte e Design  
Lead Data Scientist, Mysense.ai

dfr@esmad.ipp.pt

<https://orcid.org/0009-0001-2022-7072>

*Slides and code available at:*  
[github.com/diogoribeiro7](https://github.com/diogoribeiro7)